
Reconfigurable Acceleration of Transformer Neural Networks with Meta-Programming Strategies for Particle Physics Experiments

Author:
Filip Wojcicki

Supervisor:
Prof. Wayne Luk

Second Marker:
Prof. Alexander Tapper

June 7, 2022

Abstract

Particle Physics studies the fundamental forces and elementary particles building the Universe. In order to verify the correctness of the theories, countless experiments have to be designed and carefully executed, with the main driving force of the myriads of engineers, physicists and researchers at the Large Hadron Collider (LHC) operated by the European Organization for Nuclear Research (CERN). With the unprecedented experiments' scale comes the challenge of accurate, ultra-low latency decision-making. Transformer Neural Networks (TNN) have been proven to accomplish cutting-edge accuracy in various domains, including classification for jet tagging, which is the target of this project. However, software-centered solution implemented for CPUs and GPUs lack the inference speed needed for real-time particle triggers.

This report proposes two novel TNN-based architectures efficiently mapped to Field-Programmable Gate Arrays (FPGAs). The first one outperforms the current state-of-the-art models' GPU inference capabilities by roughly 1000 times while maintaining comparable classification accuracy. The second one trades off some of its speed for accuracy and undergoes a broad design-space exploration, which involves both pre-training and post-training quantization. The latter one leverages a custom-developed tool chain that augments existing solutions in terms of granularity and ease-of-use while following an innovative algorithm for relatively quick convergence.

In this project, several recently researched neural network components are designed to target FPGAs using High-Level Synthesis (HLS). The resulting open-sourced building blocks are both highly customizable and abstract, and aim to bridge the gap between hardware and software development, effectively reducing the time and complexity needed for creating efficient neural network hardware accelerators.

Confirm
this num-
ber

Acknowledgements

I would like to express my gratitude to Professor Wayne Luk for his guidance, insightful suggestions and constant encouragement throughout the project.

I would like to thank Professor Tapper for giving me a different view on the project's meaning and providing with the behind-the-scenes information about the LHC.

I want to thank Zhiqiang Que for his continuous technical support, our weekly meetings and always being available to answer any of my questions.

Lastly, I am very grateful for my family whose support was invaluable during this project and the degree as a whole.

Contents

List of Figures	4
List of Tables	4
1 Introduction	5
1.1 Motivation	5
1.2 Objectives and Challenges	5
1.3 Contributions and Publication	6
1.4 Report Outline	6
2 Background and related work	7
2.1 Particle Physics	7
2.1.1 Standard Model	7
2.1.2 Particle Accelerators and Triggers at LHC	8
2.1.3 Dataset and Notations	8
2.2 Machine Learning	9
2.2.1 Classification Accuracy, Area-Under-Curve, and Confusion Catrrix	9
2.2.2 Deep Neural Networks	9
2.2.3 Transformer Neural Networks and Attention	10
2.3 Reconfigurable Hardware	10
2.3.1 Landscape of Hardware for Computing	10
2.3.2 High-Level Synthesis	11
2.3.3 hls4ml Codesign Workflow	11
2.3.4 Latency, Throughput, and Hardware Resource Utilization	11
2.3.5 Serial, Parallel, and Pipelined architectures	12
2.3.6 Pareto Front and Roofline Model	13
2.4 Ethical Considerations	14
3 Models Implementation	16
3.1 Baseline Software Model	16
3.2 Ultra-Low Latency Model	16
3.2.1 Simplification and Tuning	16
3.2.2 Hardware Mapping	16
3.3 Accuracy-Focused Model	16
3.3.1 Hardware Mapping	16
3.4 Parameter Extraction for Custom Hardware	16
4 Design Space Exploration	17
4.1 Pre-training Quantization	17
4.2 Post-training Quantization	17
4.3 High-Level-Synthesis Optimization	17
5 Implementation	18
5.1 Adaptation of the PyTorch ConstituentNet architecture	18
5.2 Parameter Extraction Tool with Normalization Embedding	18
5.3 Implementation of ConstituentNet in Vivado HLS	19

5.4	Research into the <i>hls4ml</i> Library and Integration Potential	19
6	Evaluation	20
6.1	Quantitative results	20
6.2	Qualitative Results	20
6.3	Quantization Results	21
7	Conclusion	22
7.1	Future Work	22
	References	23
	Appendices	27
A	Something	28

List of Figures

2.1	Representation of different decay processes, based on the number of resulting jet clusters	8
2.2	High-level overview of the ConstituentNet architecture	10
2.3	Diagram comparing serial and parallel configurations as well as showcasing designs with and without pipelining	12
2.4	Example graph with designs plotted against quantities A/B, Pareto front highlighted	13
2.5	Example graph with computational and memory bandwidth limitations showcasing the Roofline model	14

List of Tables

Chapter 1

Introduction

1.1 Motivation

LHC is the world’s highest-energy particle collider that is capable of producing and detecting the heaviest types of particles that emerge from events such as proton-proton collisions. The detection is a challenging process as some particles, like quarks and gluons, cannot exist on their own, and they nearly instantly combine which results in collimated sprays of composite particles (hadrons) that are referred to as jets [1]. The initial particles created upon collision and their behaviors are of main interest of the physicists, which leads to jet tagging - the challenge of associating particle jets with their origin.

There are many detector types used for the analysis of the particle collisions, each based on a different physical phenomenon, which results in availability of both higher and lower level features from each event. The former have been successfully used in the past using more physically motivated machine learning (ML) algorithms, e.g. using computer vision [2]. However, more recently, various deep learning approaches have proven to outperform their predecessors [3]. It has also been found that all the detected features carry the same underlying information, with convolutional neural networks trained on higher-level data achieving nearly identical accuracy as dense neural networks trained on the data from the other end of the spectrum [4].

The information throughput of Petabytes per second collected by the LHC detectors outclasses the real-time inference capabilities of the typical state-of-the-art solutions. The real-time decision-making is often of utmost interest, hence this paper is motivated by this challenge which includes exploring various types of neural network architecture as well as the necessary infrastructure and deployment processes. Recently, *hls4ml* codesign workflow have been successfully adopted in particle physics experiments [5], which allows ML researchers and physicists to easily deploy their solutions trained using common ML frameworks on reconfigurable or application specific hardware, vastly improving the detection algorithms’ inference capabilities. However, *hls4ml* lacks support for a number of neural network architectures that have been proven to outperform the previous state-of-the-art, including graph neural networks [6, 7] and transformer neural networks [8].

1.2 Objectives and Challenges

The purpose of this project is to develop novel, hardware-aware neural network architectures as well as to establish efficient ways mapping them onto FPGAs. Another objective is to use metaprogramming strategies to integrate them into the *hls4ml* library or standalone tools, with various optimizations approaches that offer trade-offs between latency, throughput and hardware resources usage. Hence, there is an emphasis on creating parametrizable and reusable designs that can support creation of ultra-low latency systems, effectively transforming proof-of-concept implementations into optimized hardware accelerators.

The two main challenges of the project involve:

- Developing deep, complex and fast neural network models which require much lower abstraction levels than a typical ML framework. It is also crucial to stay aware of the underlying hardware architecture to exploit its strengths while keeping compile-time and run-time configuration easily accessible.
- Bridging the abstraction gap for the translation between high-level representation of neural networks and their optimized mapping to hardware. The design space exploration is a long and difficult process, which needs careful examination and analytical performance models to find the optimal solutions.

1.3 Contributions and Publication

The project aims to benefit the open-source community of ML researches that are in need of faster and more parametrizable neural network inference. The main audience for that operation are particle physicists, nonetheless, the hope is for the work to positively contribute many ML fields by both offering a reliable tool for acceleration of existing designs and providing a useful resource for learning about the nature of reconfigurable hardware and its optimization potential.

The bulk of the work and analysis conducted in this project was summarized in the paper "Accelerating Transformer Neural Networks on FPGAs for High Energy Physics Experiments" and submitted in the long paper category to the *18th International Symposium on Applied Reconfigurable Computing*. A journal article derived from this project is being prepared for publication.

1.4 Report Outline

This report begins by discussing the necessary particle physics background to understand the scope of the work, followed by the related work in the field of machine learning, with an emphasis on the state-of-the-art architectures, and a deeper dive into the reconfigurable hardware technology in chapter 2. Then in chapter 3, the two proposed novel neural network architectures are described in details, including the necessary training and processing steps. After that, chapter 4 covers both the existing and custom ways of conducting design space exploration and how they were applied in this project. Chapter 6 discusses the evaluation metrics and collected experimental results, which is concluded by chapter 7, which also proposes future work derived from this analysis.

Chapter 2

Background and related work

This chapter provides a closer look at the concepts required to understand this work. The following sections firstly discuss background and related work for topics in particle physics, then machine learning and finally reconfigurable hardware research. At the end of this chapter, the ethical issues that could arise from the project are listed and discussed.

2.1 Particle Physics

To be able to understand the scope of the project and the applicability of the work in modern research, this chapter gives an overview of the key concepts from particle physics that appear through the paper. The explanation is written for readers with no prior background in physics.

2.1.1 Standard Model

the Standard Model is a theory that describes the connections between weak, strong and electromagnetic interactions, which are three of the fundamental forces. The possible unification with the forth one - gravity - is an ongoing research [9], and while certainly outside the scope of this project, it should be noted that some of the physical experiments that this work explores aim to help with it [10, 11].

The Standard Model also provides a classification of all the elementary particles. A non-exhaustive list of them is described below, with particles that this report is concerned about (as they appear in the proton-proton collisions) being highlighted.

- Fermions
 - Leptons - participate in electroweak interactions; include electron (e^-)
 - Quarks - participate in strong interactions; include **light quarks (q)**¹ and **top (t) quark**
- Bosons
 - Gauge bosons - force carriers; include photon (γ), **W boson (W^+ , W^-)**, **Z boson**, gluons
 - Scalar bosons - give rise to mass; include **Higgs boson (H^0)**

The information about the following decay processes form the dataset of this report, with visualization in figure 2.1 (obtained from [6]). It is important to note that where applicable, the particles on the left-hand side of the arrows undergo a series of decays before reaching the right-hand side, when the only particles left are those composed of quarks and antiquarks (denoted by the vertical bar), referred to as hadrons.

$$\begin{aligned} & q/g \\ & H^0/W/Z \rightarrow q\bar{q} \\ & t \rightarrow Wq \rightarrow q\bar{q}q \end{aligned}$$

¹Light flavor quarks: up (u), down (d), charm (c), and strange (s) quarks




Figure 2.1: Representation of different decay processes, based on the number of resulting jet clusters

2.1.2 Particle Accelerators and Triggers at LHC

The two LHC experiments that are of most concern in this report are CMS and ATLAS. They are both large general-purpose particle detectors, that were notably involved in the discovery of the Higgs boson [12]. Several processing steps happen between particles colliding and theories being proven, however real-time particle detectors comprises the very first elements of this pipeline. They are composed of triggers split into the following levels [13, p.16]:

- **Level 1 trigger (L1T)** - it is implemented in hardware (FPGAs) and firmware, it is pipelined (a term explained in details in subsection 2.3.5) and it cannot allow for any dead-time, which means that it has to continuously process data with a fixed latency.
- **Level 2 trigger (L2T)** - it is implemented in hardware and software and can include regional processing.
- **Level 3 trigger (L3T)** - it is implemented in software, using farms of CPUs. It is close in behavior to non-real-time algorithms.

LHC operates in intertwined periods of operation and shutdown. The latter come from the demanding nature of the experiments that necessitates maintenance and upgrades to the apparatus and machinery, as well as the science and engineering advancements which allow for more efficient algorithms and technologies to be adopted. Very recently, after four years of break, LHC restarted experiments, which marks the beginning of "Run 3" [14]. Since their origin, the L2T and L3T have been merged into High Level Trigger (HLT) [15, p.47], which is planned to rely on thousands of multithreaded CPUs and GPUs. As for the L1T key specifications that will be used to evaluate the design in this paper, its input data frequency is 40 MHz, which with a pipeline depth of 500 results in a $12.5 \mu s$ latency, and its output frequency to HLT is equal to 750 kHz.

2.1.3 Dataset and Notations

The datasets used in this work has been simulated to mimic the 13 TeV proton-proton collisions performed at LHC, and it includes information about the most energetic jets [16] (30 [17], 50 [18], 100 [19] and 150 [20]) that were constructed using the anti- K_t clustering algorithm [21]. A number of jet representations are available in the dataset:

- High level features (HLF), which are physically inspired,
- Images, which are related to an energy heat-map,
- Constituent list, which contains jets' constituent hadrons from the following list: light quarks, top quarks, W bosons, Z bosons, and gluons.

Compared to the other two, the constituent list is a lower-level representation, however, as mentioned in chapter 1, this should not affect the classification accuracy [4]. It is also worth mentioning that a simpler dataset that contains only the HLF jet representation [22] is also used in this project as it vastly reduces the complexity of a design while offering comparable accuracy. A more thorough discussion between the differences in their use cases is carried in chapter 6.

To facilitate further analysis, this subsection also explains the notation used for the dataset as it is important throughout the report. Each constituent element \mathbf{x}^l is a 16-dimensional vector, where l denotes the index in the list:

$$\mathbf{x}^l = [x_0^l \ x_1^l \ \dots \ x_{15}^l]^T \in \mathbb{R}^{16} \quad (2.1)$$

The physical meaning of each element's dimension is not taken into consideration, and all of them are treated as equally important. A constituent list \mathbf{x}^i acts a single sample, with index i within a dataset, and it varies in terms of the number of constituents L , but it has no more of them than the dataset name suggests:

$$\hat{\mathbf{x}}^i = [\mathbf{x}^{i,0} \ \mathbf{x}^{i,1} \ \dots \ \mathbf{x}^{i,L-1}] \in \mathbb{R}^{L \times 16} \quad (2.2)$$

Using the 30 jet dataset as an example, each sample has between 1 and 30 constituents, although in the majority of samples it is the upper boundary. Hence, the whole dataset with N samples can be represented as D :

$$D = \begin{bmatrix} \hat{\mathbf{x}}^0 \\ \hat{\mathbf{x}}^1 \\ \vdots \\ \hat{\mathbf{x}}^{N-1} \end{bmatrix} = \begin{bmatrix} \mathbf{x}^{0,0} & \mathbf{x}^{0,1} & \dots & \mathbf{x}^{0,L-1} \\ \mathbf{x}^{1,0} & \mathbf{x}^{1,1} & \dots & \mathbf{x}^{1,L-1} \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{x}^{N-1,0} & \mathbf{x}^{N-1,1} & \dots & \mathbf{x}^{N-1,L-1} \end{bmatrix} \in \mathbb{R}^{N \times L \times 16} \quad (2.3)$$

The main jet datasets contain 880,000 samples regardless of the number of jets per sample, while the HLF dataset contains 830,000 examples, which are split into training and test samples in 70 : 30 and 80 : 20 proportions accordingly.

2.2 Machine Learning

The already mentioned neural networks belong to a wider field of machine learning (ML) - the study of using experience to improve algorithms. This section assumes a basic understanding of ML and gives a brief overview of the topics needed to understand the scope of the project. It then explains in more details the background and related work for the architectures involved in this research.

2.2.1 Classification Accuracy, Area-Under-Curve, and Confusion Catrrix

There are several key metrics used for assessing the success of an ML algorithm, and the following will be used throughout the report:

- **Classification accuracy** - a simple measure of the percentage of correctly classified samples.
- **Confusion matrix** - a tabular metric that compares the actual samples' classes with the predicted ones, effectively categorizing results into four groups: true positive, false positive, false negative, and true negative. This allows for an easy calculation of precision and recall values.
- **Area Under the Curve (AUC) for the Receiver Operator Characteristic (ROC)** - a more complex measure of the model's ability to correctly distinguish between classes. It can be used similarly to the classification accuracy, but it favors discriminative over representative models.

2.2.2 Deep Neural Networks

While there exist a number of ML techniques that have proven successful for various use cases at LHC, like Support Vector Machines [23] or Boosted Decision Trees [24], in the last years deep neural networks (DNN) have been proposed with improved results for applications like infrastructure monitoring [25], offline data analysis [26], and the main interest of this report - detectors' trigger mechanisms.

In many uses cases the neural networks architectures are optimized and accelerated to shorten the training time (often measured in hours) to reduce the time needed for evaluating different

design configurations and easily perform the hyperparameter search. However, this work focuses on accelerating the inference to match the extremely low latency required in the LHC detectors' L1 triggers. Although often measured in milliseconds, sub-milliseconds inference time has been achieved for this application with the use of FPGAs using architectures for basic DNN [22], and recently sub-microseconds latency for graph neural networks (GNN) [27, 28]. These implementations serve as a baseline latency for this project which aims to achieve comparable performance with higher AUC value. A commonality between the recent best performing designs is the use of the *hls4ml* codesign workflow that was mentioned in [section 1.1](#).

2.2.3 Transformer Neural Networks and Attention

A promising architecture that has been chosen as the first implementation for this project is the transformer neural network, which is a type of recurrent neural network (RNN). A recent implementation [8] called ConstituentNet outperforms previous state-of-the-art graph neural networks (GNN) implementations like JEDI-net [6] using a version of the attention mechanism [29], called self-attention. A diagram with a high-level view of the ConstituentNet can be seen in [figure 2.2](#).

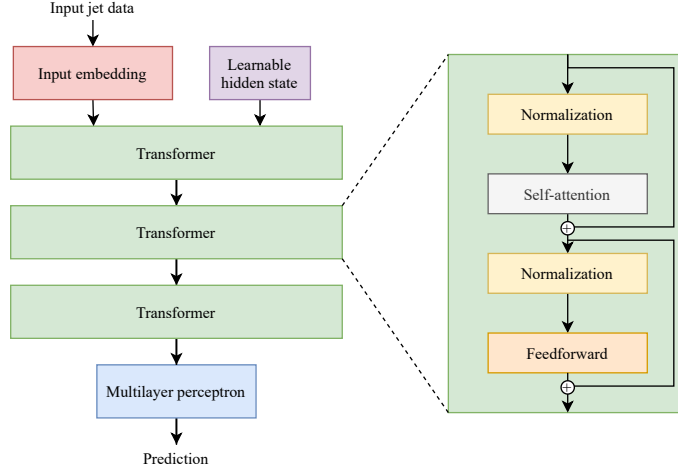


Figure 2.2: High-level overview of the ConstituentNet architecture

Given the strong results of its software implementation, an FPGA-mapped design can become an improvement over existing designs for L1Ts at LHC. A number of techniques are planned to ensure the design matches the required performance, which includes quantization and pruning [30] and careful exploration of the trade-off between latency and hardware resource utilization, which is covered in [subsection 2.3.4](#).

2.3 Reconfigurable Hardware

A significant portion of the project's work involves exploiting reconfigurable hardware to vastly reduce the inference time of state-of-the-art neural networks. This section explains in more detail the technology and characteristics of the Field-Programmable Gate Arrays (FPGA).

2.3.1 Landscape of Hardware for Computing

The modern landscape of digital integrated circuits (IC) is very rich can be divided into numerous categories depending on the technology used and expected functionality [31]. A list of platform types is described below, with the emphasis of their suitability for neural networks applications.

- **Central Processing Units (CPU)** - the most commonly found ICs that are at the core of personal computers, laptops and handheld devices. They are capable of executing a broad range of predefined instructions. As CPUs have become widely adopted in research long before the emergence of the other technologies from this list, they were the first platforms for the training and inference of neural networks with promising results back in the 1980s and 1990s for applications like high energy physics [32] or biology [33]. Although possible

to achieve speed-ups of over 10x the baseline performance with careful optimizations [34], CPUs are now consistently outperformed by more suitable technologies.

- **Graphic Processors (GPU)** - ICs specialized in graphics processing intended for displaying images. Since their original use case, due to the type of calculations involving matrix and vector operations, other applications related to cryptography and neural networks have also adopted GPUs as their main resource. In the former domain, cryptocurrency mining has transitioned from CPU to GPU to increase profitability [35], while for the latter, the more powerful hardware drastically reduced training and inference times, thus allowing for deeper and more complex architectures yielding higher accuracy [36, 37].
- **Application Specific Integrated Circuits (ASIC)** - as suggested by the name, those are the custom designed ICs heavily specialized for a particular use. It is hard to generalize them, as the use cases can cover any modern computing problem, but the commonality is a vast improvement in performance and power usage compared to more general purpose solutions. However, the long and expensive development process pose an extremely high barrier to entry for most users. Fortunately, off-the-shelf products like the Graphcore Intelligence Processing Units [38] that are designed specifically with machine learning applications in mind as well as other custom designs [39, 40] are starting to offer a compelling platform for working with neural networks.
- **Field-Programmable Gate Arrays (FPGA)** - differently from the previous listed IC types, FPGAs are not manufactured for a specific use case, and in fact, they can be reprogrammed to be a platform for a different application at any time. The reprogrammability comes at a cost of performance and power consumption compared to ASICs [41], but at the same time outperforms GPUs in these regards [42, 43]. It is also suggested, that with some technological improvements focused on ML applications, FPGAs can narrow the gap between ASICs without needing to stick to one particular design [41, 44, 45].

FPGAs offer an interesting trade-off between implementation effort and acceleration potential when it comes to neural networks and for that reason they have been chosen the target technology in this report. The following subsections give a closer look at some of their characteristics.

2.3.2 High-Level Synthesis

For many years, FPGAs have been modelled using register-transfer level (RTL) design abstraction with the use of hardware description languages like Verilog or VHDL. However, to increase productivity and allow for a more convenient design state space exploration, a more abstract modelling process called High-Level Synthesis (HLS) can be adopted. The design can be expressed in a software programming language like C or C++, which are automatically optimized and transformed to an equivalent RTL. This is especially beneficial in research, where compared to industrial environment, it is more likely that designers can afford slightly lower quality of results for increased productivity. In fact, a recent study shows that on average, only one third of design time and half of the lines of code are needed for an equivalent project done in HLS in comparison to RTL while the quality of results varies and can even outperform the RTL implementations [46].

This report's work is based on Xilinx Vivado HLS design suite. When developing a solution, it is important to note, that the synthesis process can take a significant amount of time (a couple of hours on a modern powerful computer), and so there exist two simulation methods - a C-simulation that can quickly and directly evaluate a C/C++ benchmark against the software implementation of the design, and a more truthful, cosimulation that firstly synthesizes the design and the test bench to RTL and then performs an RTL simulation. A final, definitive evaluation of the results requires programming a target FPGA with the generated bit stream of the design and exchanging input/output data with a program that usually runs on a CPU.

2.3.3 hls4ml Codesign Workflow

2.3.4 Latency, Throughput, and Hardware Resource Utilization

To understand the differences between hardware designs targeted at similar functionality, it is worth considering the following characteristics:

- **Latency** - A time measure of a system between receiving an input signal and producing a *corresponding* output. It is crucial in real-time processing where it has to be lower than the period between subsequent input samples. Depending on the application, latency in the microseconds range can be expected from an FPGA.
- **Throughput** - A rate of samples processed in a unit of time. For architectures that only start to process new elements after the previous one has finished, it is equal to latency. However, in modern ICs, especially in FPGAs, it is one of the defining metrics of performance and designs tend to exploit pipelining and parallelizability to marginally trade off their latency to increase it.
- **Resource utilization** - A more complicated, often multidimensional, metric that describes the raw number or ratio of total usage of the hardware components on an FPGA. Typically, the higher it becomes, the more power is drawn by an FPGA, however, it is most often used to guide the design process to avoid running out of a certain resource and potentially deploy an alternative method that can be implemented using a different, less contested resource.

To fully understand the trade-offs between designs, one cannot forget about the metric related to the specific task that is accelerated in hardware. In the case of this report, classification accuracy and AUC described in [subsection 2.2.1](#), will also play key roles in evaluating various configurations.

2.3.5 Serial, Parallel, and Pipelined architectures

Hardware architectures use components that can be configured in different ways depending on the overall goal or a limiting factor. The high-level configurations are displayed in figure 2.3 and can be described as follows:

- **Serial** - elements are arranged in a chain, processing one after another. This way uses less resources than an equivalent parallel configuration.
- **Parallel** - elements share a common input and start processing data at the same time. This way ends in a lower latency than an equivalent serial configuration.
- **Pipelined** - a more sophisticated arrangement, in which subsequent processing blocks (that can be either placed serially or in parallel) form a pipeline of processing stations separated by simple storage elements called pipeline registers. This maximizes the usage of the design blocks, hence increasing throughput with a minimal sacrifice of latency and resource usage.

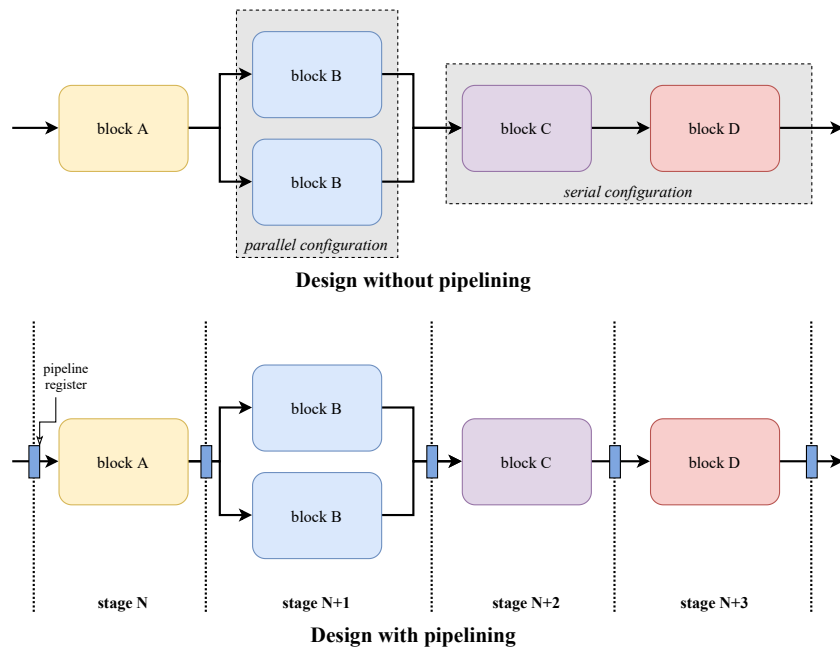


Figure 2.3: Diagram comparing serial and parallel configurations as well as showcasing designs with and without pipelining

2.3.6 Pareto Front and Roofline Model

To make an informed design decision, various architectures can be compared by arranging them on a dependency graph (e.g. latency vs resource usage) and observing the Pareto front - the set of solutions for which there are no better ones in regard to one quality given that the other measure is not worse. The slightly complex definition can be easily understood from figure 2.4, which also highlights another use of this method - finding design configurations that are yet to be explored.




Figure 2.4: Example graph with designs plotted against quantities A/B, Pareto front highlighted

Another intuitive performance visualization comes in the form of the Roofline model, which compares the obtained results with theoretical limits coming from inherent hardware limitations like clock frequency or memory bandwidth. An example can be seen in fig 2.5.




Figure 2.5: Example graph with computational and memory bandwidth limitations showcasing the Roofline model

some graphics to potentially add: fpga lattice, hls to rtl flow, rtl to bit stream flow

difficulty: rtl > hls > python hl4ml, draw comparison with assembly

FPGA are very hard-coded -> make the code deployable on any platform with optimal settings automatically

HLS is difficult, so coding hardware in Python is desired -> make it easy for engineers and physicists to design systems

Metaprogramming allows for optimizations and customisability

2.4 Ethical Considerations

The purpose of this project is to advance the next-generation particle physics experiments. There are two main aspects that need to be considered - the development of a hardware-mapped transformer neural network architecture and the easy-to-access translation and optimization toolchain for efficiently expressing networks in common machine learning frameworks.

The first feature is aimed at a purely civilian, scientific audience and it is tailored towards particle collision datasets. With that in mind, it is important to mention that, as with most machine learning research, there is potential for a misuse of the acceleration techniques towards a military or malevolent application that could negatively impact the society. However, this also means that there is a low risk for new emerging threats; rather the already present ones could become more serious. Fortunately, this should result in existing harm prevention measures staying intact or solely requiring adjustments to their accuracy or speed thresholds.

With the second element's goal of making the creation and deployment of neural networks more accessible, it could be argued that this may in turn increase the number of physics experiments requiring high energy consumption, like those at LHC [47], thus negatively affecting the environment. However, this is considered a very low likely cause of action, as the research work of this project is aimed at helping already running experiments and more importantly, the negative environmental

implications (for which there are various mitigation strategies [48, 49]) are heavily outweighed by potential beneficial technological advancements coming from the scientific discoveries.

Despite the aforementioned ethical issues, the project is aimed at benefitting the open-source scientific community world-wide. Its outcome could lead to a much more accessible and efficient inference methods that are applicable in many domains outside particle physics.

Chapter 3

Models Implementation

3.1 Baseline Software Model

3.2 Ultra-Low Latency Model

3.2.1 Simplification and Tuning

3.2.2 Hardware Mapping

3.3 Accuracy-Focused Model

3.3.1 Hardware Mapping

3.4 Parameter Extraction for Custom Hardware

tool for extracting weight and biases

In order to generate files with weights and biases that are required for initializing the memory on an FPGA, a tool was developed that takes a PyTorch pre-trained model, extracts all the information, and splits them accordingly with the required format. What is more, layers responsible for normalization can be chosen to have their mean and variance calculation embedded into weights and biases to significantly reduce the processing required on an FPGA by omitting the division and square root operations. The mathematical derivation of this approach starts with the batch norm formula:

$$y = \frac{x - E[x]}{\sqrt{Var[x] + \epsilon}} * \gamma + \beta$$

The expected value and variance are treated as learnable parameter of a dataset and are extracted after the training has been completed. Hence, the calculation becomes:

$$y = x * \frac{\gamma}{\sqrt{Var + \epsilon}} + \beta - \frac{\gamma * E}{\sqrt{Var + \epsilon}} = x * W + b$$

The newly calculated values for W and b represent the updated weights and biases of the normalization layer, that can be then implemented in hardware in a much simpler way. Independently of the implementation in this work, a similar idea has been proposed and successfully used as an optimization in the past [50].

tool for embedding norm stats for layer norm as running stats not collected

Chapter 4

Design Space Exploration

4.1 Pre-training Quantization

PyTorch Eager Mode

PyTorch FX Graph Mode

Brevitas

QPyTorch

4.2 Post-training Quantization

Custom tool

4.3 High-Level-Synthesis Optimization

ScaleHLS

MLIR

Chapter 5

Implementation

As depicted in the figure ??, part of the project plan from ?? has already been implemented as of the time of publishing this report. This was done thanks to the smaller workload of the Autumn term in comparison to the Spring term as well as the significant effort over the Winter break. This 'head start' is hoped to allow for a deeper state space exploration and a more refined final architecture and in case of faster than expected working pace, further extensions related to the *hls4ml* library and automatic optimizations will also be considered. The accomplishments so far can be categorized into four domains covered in the following sections.

5.1 Adaptation of the PyTorch ConstituentNet architecture

Thanks to the existing code base with an implementation, it was easier to understand the smaller details that were not fully explained in the original paper [8]. However, many aspects of the provided code served as proof-of-concept and are suspected to have been changed after the publishing, as a new model could not have had been trained, nor the provided one could have had been evaluated. Without the help of the original author, a severe investigation and fixing process were required to progress the software implementation into a usable state. Despite those difficulties, the time was well spent on finding potential optimization points for the later stage of the project. Moreover, frequent reporting hooks were added in between the existing network layers, which allows for generating an inner view of the calculations happening on the CPU that gives the opportunity for direct, step-by-step comparison with the HLS implementation. Ultimately, the code base has reached a state where it is convenient to train models with different parameters and evaluate them against the datasets.

5.2 Parameter Extraction Tool with Normalization Embedding

In order to generate files with weights and biases that are required for initializing the memory on an FPGA, a tool was developed that takes a PyTorch pre-trained model, extracts all the information, and splits them accordingly with the required format. What is more, layers responsible for normalization can be chosen to have their mean and variance calculation embedded into weights and biases to significantly reduce the processing required on an FPGA by omitting the division and square root operations. The mathematical derivation of this approach starts with the batch norm formula:

$$y = \frac{x - E[x]}{\sqrt{Var[x] + \epsilon}} * \gamma + \beta$$

The expected value and variance are treated as learnable parameter of a dataset and are extracted after the training has been completed. Hence, the calculation becomes:

$$y = x * \frac{\gamma}{\sqrt{Var + \epsilon}} + \beta - \frac{\gamma * E}{\sqrt{Var + \epsilon}} = x * W + b$$

The newly calculated values for W and b represent the updated weights and biases of the normalization layer, that can be then implemented in hardware in a much simpler way. Independently of the implementation in this work, a similar idea has been proposed and successfully used as an optimization in the past [50].

5.3 Implementation of ConstituentNet in Vivado HLS

An effort has been made to express the original version of ConstituentNet in HLS based off an empty skeleton code generated by *hls4ml*. So far, the architecture, including the blocks responsible for the transformer and the self-attention calculations, have been designed and connected accordingly. Several layer types, namely the densely-connected, normalization and activation ones have already been successfully validated. A variant of the densely-connected layer used for matrix-matrix multiplication, along with a few smaller layers, are yet to be validated, at which point, the architecture would be ready for an evaluation against the PyTorch implementation.

5.4 Research into the *hls4ml* Library and Integration Potential

The current HLS implementation benefits from a number of existing layers included as part of the *hls4ml* library. This has been done as a way to shorten the time needed for the initial HLS implementation that will serve as the baseline for further optimizations as well as to research the structure and methods used in the library to allow for a smoother integration into the codesign workflow once the initial objective is met.

Chapter 6

Evaluation

This chapter outlines the proposed evaluation plan for the project. The first objective of developing and optimizing a state-of-the-art neural network in hardware can be evaluated quantitatively, while integrating it into the *hls4ml* library and making it easy for new users to use requires a more qualitative approach.

Analytical models for latency/resources?

6.1 Quantitative results

The following describes the quantities to be measured for each neural network design:

- Classification accuracy, AUC and confusion matrix on a validation dataset
- Inference latency and throughput when running on the target platform
- Hardware resource utilization (exact values for comparison with other platforms and percentage of available resources for understanding limitations):
 - Block RAM (BRAM) and Ultra RAM (URAM)
 - Digital Signal Processing units (DSP)
 - Flip-Flops (FF)
 - Look-Up Tables (LUT)

In the early stages of the project, the above quantities will be measured from the results from the simulation and synthesis reports. At a later stage, the best designs will be run on actual hardware platforms to validate them under real-life use cases. The platform planned for this part is an Intel Stratix V FPGA hosted in a Maxeler MPC-X dataflow node with 8 Maia dataflow engines and 48 GB of DRAM. A consideration is also planned for the specific hardware used in the LHC L1T detectors and its available resources, which although cannot be directly tested on, can guide the state space exploration.

Apart from clear design improvements, it is predicted that most evaluated designs will offer trade-offs between classification accuracy, AUC, inference throughput and hardware utilization. It is not possible to find a design that is superior in every way, hence a Pareto front and the Roofline model will play the key roles in understanding the overall performance and selecting configuration with specific needs in mind.

6.2 Qualitative Results

To assess the success of enhancing the *hls4ml* library, qualitative comparisons will be drawn between it and the already existing neural network components and architectures. Depending on the project's timeline, it is possible that the improvements can get official approval and get merged

into the main repository, however if this is not feasible before the final deadline, current users of the library will be surveyed and their opinion will be taken into consideration instead.

6.3 Quantization Results

pre-training quantization compared to varying floating-point widths

float16 doesnt learn anything (acc 20%) as its range is too small and we cannot consider normalizing inputs coz its real time system

brevitas only gets 34% accuracy

pytorch quantization is too experimental and doesnt support the model

post-training quantization

somewhere: fuse batch norm to linear???

Chapter 7

Conclusion

7.1 Future Work

bullet points

References

- [1] CERN. Jets at CMS and the determination of their energy scale | CMS experiment, .
- [2] Josh Cogan, Michael Kagan, Emanuel Strauss, and Ariel Schwartzman. Jet-images: computer vision inspired techniques for jet tagging. *The journal of high energy physics*, 2015(2):1–16, Feb 18, 2015. doi: 10.1007/JHEP02(2015)118. URL [https://link.springer.com/article/10.1007/JHEP02\(2015\)118](https://link.springer.com/article/10.1007/JHEP02(2015)118).
- [3] Luke de Oliveira, Michael Kagan, Lester Mackey, Benjamin Nachman, and Ariel Schwartzman. Jet-images — deep learning edition. *The journal of high energy physics*, 2016(7):1–32, Jul 13, 2016. doi: 10.1007/JHEP07(2016)069. URL [https://link.springer.com/article/10.1007/JHEP07\(2016\)069](https://link.springer.com/article/10.1007/JHEP07(2016)069).
- [4] Liam Moore, Karl Nordstrom, Sreedevi Varma, and Malcolm Fairbairn. Reports of my demise are greatly exaggerated: n -subjettiness taggers take on jet images. *SciPost physics*, 7(3):036, Sep 24, 2019. doi: 10.21468/SciPostPhys.7.3.036. URL <https://hal.archives-ouvertes.fr/hal-01851157>.
- [5] Farah Fahim, Benjamin Hawks, Christian Herwig, James Hirschauer, Sergio Jindariani, Nhan Tran, Luca P. Carloni, Giuseppe Di Guglielmo, Philip Harris, Jeffrey Krupa, Dylan Rankin, Manuel Blanco Valentin, Josiah Hester, Yingyi Luo, John Mamish, Seda Orgrenci-Memik, Thea Aarrestad, Hamza Javed, Vladimir Loncar, Maurizio Pierini, Adrian Alan Pol, Sioni Summers, Javier Duarte, Scott Hauck, Shih-Chieh Hsu, Jennifer Ngadiuba, Mia Liu, Duc Hoang, Edward Kreinar, and Zhenbin Wu. hls4ml: An open-source codesign workflow to empower scientific low-power machine learning devices. Mar 9, 2021. URL <https://arxiv.org/abs/2103.05579>.
- [6] Harvey B. Newman, Avikar Periwal, Maria Spiropulu, Javier M. Duarte, Maurizio Pierini, Eric A. Moreno, Aidana Serikova, Olmo Cerri, Jean-Roch Vlimant, and Thong Q. Nguyen. JEDI-net: a jet identification algorithm based on interaction networks. *The European physical journal. C, Particles and fields*, 80(1):1–15, Aug 14, 2019. doi: 10.1140/epjc/s10052-020-7608-4. URL <http://cds.cern.ch/record/2688535>.
- [7] Abdelrahman Elabd, Vesal Razavimaleki, Shi-Yu Huang, Javier Duarte, Markus Atkinson, Gage DeZoort, Peter Elmer, Jin-Xuan Hu, Shih-Chieh Hsu, Bo-Cheng Lai, Mark Neubauer, Isobel Ojalvo, and Savannah Thais. Graph neural networks for charged particle tracking on FPGAs. Dec 3, 2021. URL <https://arxiv.org/abs/2112.02048>.
- [8] Xinyang Yuan. Constituentnet: Learn to solve jet tagging through attention. Technical report, -09-22 2021.
- [9] K. Krasnov and R. Percacci. Gravity and unification: a review. *Classical and quantum gravity*, 35(14):143001, Jun 14, 2018. doi: 10.1088/1361-6382/aac58d. URL <https://iopscience.iop.org/article/10.1088/1361-6382/aac58d>.
- [10] J. Walz, P. Grandemange, B. Vallery, P. Debu, Y. Matsuda, N. Sillitoe, M. Charlton, T. Mortensen, A. Husson, F. Nez, Y. Sacquin, F. Biraben, G. Dufour, Y. Yamazaki, O. Dalkarov, P. A. Hervieux, J. M. Rey, R. Guerout, L. Liskay, S. Wronka, D. Lunney, S. Wolf, A. Voronin, V. Nesvizhevsky, M. Valdes, D. Banerjee, N. Madsen, N. Kuroda, P. Froelich, K. Khabarova, L. Hilico, P. Comini, C. Regenfus, S. Reynaud, S. Jonsell, A. Rubbia, G. Manfredi, J. M. Heinrich, J. P. Karr, C. I. Szabo-Foster, S. Eriksson, M. Staszczak, H. Torii,

- P. Indelicato, Y. Nagashima, P. Dupre, P. Clade, A. Lambrecht, A. Mohri, J. M. Raymond, D. P. Werf, B. Mansoulie, S. Guellati, A. Douillet, P. Crivelli, A. M. M. Leite, P. Perez, D. Brook-Roberge, F. Schmidt-Kaler, and N. Kolachevsky. The GBAR antimatter gravity experiment. Technical Report 233, Springer International Publishing, 2015. URL <http://cds.cern.ch/record/2055685>.
- [11] D. Pagano, M. Caccia, J. Fesel, S. Gerber, C. Malbrunot, S. R. Muller, G. Cerchiari, M. Giannamarchi, F. Guatieri, L. Smestad, S. Aghion, E. Widmann, F. Prelz, N. Zurlo, F. Sorrentino, H. Sandaker, J. Robert, A. Evans, S. Haider, P. Yzombard, A. Gligorova, L. Di Noto, P. Lansonneur, O. Khalidova, D. Krasnicky, R. Santoro, P. Lebrun, M. Fani, A. Kellerbauer, Z. Mazzotta, H. Holmestad, V. Petracek, S. Mariazzi, L. Penasa, D. Comparat, M. Prevedelli, G. Nebbia, J. Zmeskal, A. Hinterberger, C. Zimmer, I. C. Tietje, R. Caravita, C. Amsler, N. Pacifico, G. Testera, L. Ravelli, M. Doser, R. Ferragut, A. Rotondi, M. Oberthaler, B. Rienaecker, G. Bonomi, A. Demetrio, P. Nedelec, O. M. Rohne, A. Fontana, R. S. Brusa, F. Castelli, J. Marton, V. Matveev, G. Consolati, and V. Lagomarsino. Gravity and antimatter: the AEgIS experiment at CERN. Technical Report 1342, IOP Publishing, 2020. URL <http://cds.cern.ch/record/2714100>.
- [12] Brian Greene. How the higgs boson was found, July 2013. URL <https://www.smithsonianmag.com/science-nature/how-the-higgs-boson-was-found-4723520/>.
- [13] Trigger, DAQ and FPGAs. URL <http://www.hep.ph.imperial.ac.uk/~tapper/lecture/trigger.pdf>.
- [14] Sean Keane. CERN’s large hadron collider restarts after three-year upgrade. URL <https://www.cnet.com/science/cerns-large-hadron-collider-restarts-after-three-year-upgrade/>.
- [15] Alex Tapper. Triggering at collider experiments. URL <http://www.hep.ph.imperial.ac.uk/~tapper/lecture/CMSIndia-2020.pdf>.
- [16] E. Coleman, M. Freytsis, A. Hinzmann, M. Narain, J. Thaler, N. Tran, and C. Vernieri. The importance of calorimetry for highly-boosted jet substructure. *Journal of instrumentation*, 13(1):T01003, Jan 9, 2018. doi: 10.1088/1748-0221/13/01/T01003. URL <https://search.proquest.com/docview/2365693051>.
- [17] Maurizio Pierini, Javier Mauricio Duarte, Nhan Tran, and Marat Freytsis. HLS4ML LHC jet dataset (30 particles), . URL <https://doi.org/10.5281/zenodo.3601436>.
- [18] Maurizio Pierini, Javier Mauricio Duarte, Nhan Tran, and Marat Freytsis. HLS4ML LHC jet dataset (50 particles), . URL <https://doi.org/10.5281/zenodo.3601443>.
- [19] Maurizio Pierini, Javier Mauricio Duarte, Nhan Tran, and Marat Freytsis. HLS4ML LHC jet dataset (100 particles), . URL <https://doi.org/10.5281/zenodo.3602254>.
- [20] Maurizio Pierini, Javier Mauricio Duarte, Nhan Tran, and Marat Freytsis. HLS4ML LHC jet dataset (150 particles), . URL <https://doi.org/10.5281/zenodo.3602260>.
- [21] Matteo Cacciari, Gavin P. Salam, and Gregory Soyez. The anti-kt jet clustering algorithm. *The journal of high energy physics*, 2008:063, Apr 1, 2008. doi: 10.1088/1126-6708/2008/04/063. URL <http://iopscience.iop.org/1126-6708/2008/04/063>.
- [22] Edward Kreinar, Jennifer Ngadiuba, Zhenbin Wu, Philip Harris, Maurizio Pierini, Ryan Rivera, Song Han, Javier Duarte, Benjamin Kreis, Nhan Tran, and Sergo Jindariani. Fast inference of deep neural networks in FPGAs for particle physics. Technical Report 13, Institute of Physics (IOP), Apr 16, 2018. URL <http://cds.cern.ch/record/2316331>.
- [23] G. Valentino, R. W. Assmann, R. Bruce, and N. Sammut. Classification of LHC beam loss spikes using support vector machines. pages 355–358. IEEE, Jan 2012. doi: 10.1109/SAMI.2012.6208988. URL <https://ieeexplore.ieee.org/document/6208988>.
- [24] Tianqi Chen and Tong He. Higgs Boson Discovery with Boosted Trees. In Glen Cowan, Cecile Germain, Isabelle Guyon, Balázs Kegl, and David Rousseau, editors, *Proceedings of the NIPS 2014 Workshop on High-energy Physics and Machine Learning*, volume 42 of *Proceedings of*

- Machine Learning Research*, pages 69–80, Montreal, Canada, 13 Dec 2015. PMLR. URL <https://proceedings.mlr.press/v42/chen14.html>.
- [25] Andrzej Skoczen, Maciej Wielgosz, and Matej Mertik. Using LSTM recurrent neural networks for monitoring the LHC superconducting magnets. Technical Report 867, Elsevier B.V, Nov 18, 2016. URL <http://cds.cern.ch/record/2234465>.
 - [26] Jie Ren, Lei Wu, and Jin Min Yang. Unveiling CP property of top-higgs coupling with graph neural networks at the LHC. *Physics letters. B*, 802:135198, Mar 10, 2020. doi: 10.1016/j.physletb.2020.135198. URL <https://dx.doi.org/10.1016/j.physletb.2020.135198>.
 - [27] Edward Kreinar, Zhenbin Wu, Gianluca Cerminara, Kinga Wozniak, Gerrit Van Onsem, Marcel Rieger, Giuseppe Di Guglielmo, Jan Kieseler, Shah Rukh Qasim, Sioni Summers, Sergo Jindariani, Jennifer Ngadiuba, Mia Liu, Philip Harris, Maurizio Pierini, Vladimir Loncar, Kevin Pedro, Yutaro Iiyama, Javier Duarte, Dylan Rankin, Nhan Tran, and Abhijay Gupta. Distance-weighted graph neural networks on FPGAs for real-time particle reconstruction in high energy physics. Technical Report 3, Frontiers Media S.A, 2020. URL <http://cds.cern.ch/record/2728798>.
 - [28] Abdelrahman Elabd, Vesal Razavimaleki, Shi-Yu Huang, Javier Duarte, Markus Atkinson, Gage DeZoort, Peter Elmer, Jin-Xuan Hu, Shih-Chieh Hsu, Bo-Cheng Lai, Mark Neubauer, Isobel Ojalvo, and Savannah Thais. Graph neural networks for charged particle tracking on FPGAs. Dec 3, 2021. URL <https://arxiv.org/abs/2112.02048>.
 - [29] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. Jun 12, 2017. URL <https://arxiv.org/abs/1706.03762>.
 - [30] Tailin Liang, John Glossner, Lei Wang, Shaobo Shi, and Xiaotong Zhang. Pruning and quantization for deep neural network acceleration: A survey. *Neurocomputing (Amsterdam)*, 461:370–403, Oct 21, 2021. doi: 10.1016/j.neucom.2021.07.045. URL <https://dx.doi.org/10.1016/j.neucom.2021.07.045>.
 - [31] Mohammadreza Najafi, Kaiwen Zhang, Mohammad Sadoghi, and Hans-Arno Jacobsen. Hardware acceleration landscape for distributed real-time analytics: Virtues and limitations. pages 1938–1948. IEEE, Jun 2017. ISBN 1063-6927. doi: 10.1109/ICDCS.2017.194. URL <https://ieeexplore.ieee.org/document/7980135>.
 - [32] Dagli and Lammers. Possible applications of neural networks in manufacturing. page 605 vol.2. IEEE TAB Neural Network Committee, 1989. doi: 10.1109/IJCNN.1989.118423. URL <https://ieeexplore.ieee.org/document/118423>.
 - [33] Cathy Wu, Michael Berry, Sailaja Shivakumar, and Jerry McLarty. Neural networks for full-scale protein sequence classification: Sequence encoding with singular value decomposition. *Machine learning*, 21(1):177, Oct 1, 1995. doi: 10.1023/A:1022677900508.
 - [34] Vincent Vanhoucke, Andrew Senior, and Mark Z. Mao. Improving the speed of neural networks on CPUs. In *Deep Learning and Unsupervised Feature Learning Workshop, NIPS 2011*, 2011.
 - [35] Sainathan Ganesh Iyer and Anurag Dipakumar Pawar. GPU and CPU accelerated mining of cryptocurrencies and their financial analysis. pages 599–604. IEEE, Aug 2018. doi: 10.1109/I-SMAC.2018.8653733. URL <https://ieeexplore.ieee.org/document/8653733>.
 - [36] Gang Chen, Haitao Meng, Yucheng Liang, and Kai Huang. GPU-accelerated real-time stereo estimation with binary neural network. *IEEE transactions on parallel and distributed systems*, 31(12):2896–2907, Dec 1, 2020. doi: 10.1109/TPDS.2020.3006238. URL <https://ieeexplore.ieee.org/document/9130887>.
 - [37] Qianru Zhang, Meng Zhang, Tinghuan Chen, Zhifei Sun, Yuzhe Ma, and Bei Yu. Recent advances in convolutional neural network acceleration. *Neurocomputing (Amsterdam)*, 323: 37–51, Jan 5, 2019. doi: 10.1016/j.neucom.2018.09.038. URL <https://dx.doi.org/10.1016/j.neucom.2018.09.038>.
 - [38] Graphcore. Graphcore intelligence processing unit. URL <https://www.graphcore.ai/products/ipu>.

- [39] Phil Knag, Jung Kuk Kim, Thomas Chen, and Zhengya Zhang. A sparse coding neural network ASIC with on-chip learning for feature extraction and encoding. *IEEE journal of solid-state circuits*, 50(4):1070–1079, Apr 2015. doi: 10.1109/JSSC.2014.2386892. URL <https://ieeexplore.ieee.org/document/7015626>.
- [40] K. Venkata Ramanaiah and Cyril Prasanna Raj. ASIC implementation of neural network based image compression. *International Journal of Computer Theory and Engineering*, pages 494–498, 2011. doi: 10.7763/IJCTE.2011.V3.356.
- [41] Andrew Boutros, Sadegh Yazdanshenas, and Vaughn Betz. You cannot improve what you do not measure. *ACM transactions on reconfigurable technology and systems*, 11(3):1–23, Dec 22, 2018. doi: 10.1145/3242898. URL <http://dl.acm.org/citation.cfm?id=3242898>.
- [42] Eriko Nurvitadhi, Ganesh Venkatesh, Jaewoong Sim, Debbie Marr, Randy Huang, Jason Ong Gee Hock, Yeong Tat Liew, Krishnan Srivatsan, Duncan Moss, Suchit Subhaschandra, and Guy Boudoukh. Can FPGAs beat GPUs in accelerating next-generation deep neural networks? FPGA '17, pages 5–14. ACM, Feb 22, 2017. doi: 10.1145/3020078.3021740. URL <http://dl.acm.org/citation.cfm?id=3021740>.
- [43] Yixing Li, Zichuan Liu, Kai Xu, Hao Yu, and Fengbo Ren. A GPU-outperforming FPGA accelerator architecture for binary convolutional neural networks. *ACM journal on emerging technologies in computing systems*, 14(2):1–16, Jul 27, 2018. doi: 10.1145/3154839. URL <http://dl.acm.org/citation.cfm?id=3154839>.
- [44] Eriko Nurvitadhi, David Sheffield, Jaewoong Sim, Asit Mishra, Ganesh Venkatesh, and Debbie Marr. Accelerating binarized neural networks: Comparison of FPGA, CPU, GPU, and ASIC. pages 77–84. IEEE, Dec 2016. doi: 10.1109/FPT.2016.7929192. URL <https://ieeexplore.ieee.org/document/7929192>.
- [45] Eriko Nurvitadhi, Jaewoong Sim, David Sheffield, Asit Mishra, Srivatsan Krishnan, and Debbie Marr. Accelerating recurrent neural networks in analytics servers: Comparison of FPGA, CPU, GPU, and ASIC. pages 1–4. EPFL, Aug 2016. doi: 10.1109/FPL.2016.7577314. URL <https://ieeexplore.ieee.org/document/7577314>.
- [46] Sakari Lahti, Panu Sjoval, Jarno Vanne, and Timo D. Hamalainen. Are we there yet? a study on the state of high-level synthesis. *IEEE transactions on computer-aided design of integrated circuits and systems*, 38(5):898–911, May 2019. doi: 10.1109/TCAD.2018.2834439. URL <https://ieeexplore.ieee.org/document/8356004>.
- [47] CERN. Facts and figures about the LHC | CERN, . URL <https://home.cern/resources/faqs/facts-and-figures-about-lhc>.
- [48] R. Guida, M. Capeans, and B. Mandelli. Characterization of RPC operation with new environmental friendly mixtures for LHC application and beyond. *Journal of Instrumentation*, 11(07):C07016–C07016, jul 2016. doi: 10.1088/1748-0221/11/07/c07016. URL <https://doi.org/10.1088/1748-0221/11/07/c07016>.
- [49] M. Capeans, R. Guida, and B. Mandelli. Strategies for reducing the environmental impact of gaseous detector operation at the CERN LHC experiments. *Nuclear Instruments and Methods in Physics Research Section A: Accelerators, Spectrometers, Detectors and Associated Equipment*, 845:253–256, 2017. doi: <https://doi.org/10.1016/j.nima.2016.04.067>. URL <https://www.sciencedirect.com/science/article/pii/S0168900216302807>. ID: 271580.
- [50] Hongxiang Fan, Shuanglong Liu, Martin Ferianc, Ho-Cheung Ng, Zhiqiang Que, Shen Liu, Xinyu Niu, and Wayne Luk. A real-time object detection accelerator with compressed SSDLite on FPGA. pages 14–21. IEEE, Dec 2018. doi: 10.1109/FPT.2018.00014. URL <https://ieeexplore.ieee.org/document/8742299>.

Appendices

Appendix A

Something

something

Notes

Confirm this number	1
Possibly clear page	7
some graphics to potentially add: fpga lattice, hls to rtl flow, rtl to bit stream flow . . .	14
difficulty: rtl > hls > python hl4ml, draw comparison with assembly	14
FPGA are very hard-coded -> make the code deployable on any platform with optimal settings automatically	14
HLS is difficult, so coding hardware in Python is desired -> make it easy for engineers and physicists to design systems	14
Metaprogramming allows for optimizations and customisability	14
tool for extracting weight and biases	16
tool for embedding norm stats for layer norm as running stats not collected	16
.	17
PyTorch Eager Mode	17
PyTorch FX Graph Mode	17
Brevitas	17
QPyTorch	17
Custom tool	17
ScaleHLS	17
MLIR	17
Analytical models for latency/resources?	20
pre-training quantization compared to varying floating-point widths	21
float16 doesnt learn anything (acc 20%) as its range is too small and we cannot consider normalizing inputs coz its real time system	21
brevitas only gets 34% accuracy	21
pytorch quantization is too experimental and doesnt support the model	21
post-training quantization	21
somewhere: fuse batch norm to linear???	21
bullet points	22
something	28